

# Why cgo is **slow**

Filippo Valsorda

```
package main
```

```
// int my_C_function(int a, int b) {  
//     return a + b;  
// }
```

```
import "C"
```

```
func main() {  
    a, b := C.int(40), C.int(2)  
    c := C.my_C_function(a, b)  
    println(a, b, c)  
}
```

cgo is a FFI

(Foreign Function Interface)

# I like FFIs.

- From cgo back to Go @ GopherCon 2016

<https://speakerdeck.com/ilosottile/from-cgo-back-to-go-gophercon-2016>

- rustgo: Building your own FFI @ GothamGo 2017

<https://speakerdeck.com/ilosottile/calling-rust-from-go-without-cgo-at-gothamgo-2017>

- Why cgo is slow @ CapitalGo 2018

Hi!

|                 |   |
|-----------------|---|
| C function call | 2.364 ns  |
| Java FFI        | 9.01 ns   |
| Rust FFI        | 2.386 ns  |
| LuaJIT FFI      | 1.81 ns (!) <a href="https://nullprogram.com/blog/2018/05/27/">https://nullprogram.com/blog/2018/05/27/</a> |
| Node.js FFI     | 18.33 ns  |
| cgo             | 75.95 ns  |

<https://github.com/dyu/ffi-overhead>

| name      | old time/op     | new time/op     | delta  |
|-----------|-----------------|-----------------|--------|
| CgoCall-4 | 63.1ns $\pm$ 3% | 57.1ns $\pm$ 0% | -9.43% |

|                 |   |
|-----------------|---|
| C function call | 2.364 ns  |
| Java FFI        | 9.01 ns   |
| Rust FFI        | 2.386 ns  |
| LuaJIT FFI      | 1.81 ns (!) <a href="https://nullprogram.com/blog/2018/05/27/">https://nullprogram.com/blog/2018/05/27/</a> |
| Node.js FFI     | 18.33 ns  |
| cgo             | 68.77 ns  |

<https://github.com/dyu/ffi-overhead>

|                 |   |
|-----------------|---|
| C function call | 2.364 ns  |
| Java FFI        | 9.01 ns   |
| Rust FFI        | 2.386 ns  |
| LuaJIT FFI      | 1.81 ns (!) <a href="https://nullprogram.com/blog/2018/05/27/">https://nullprogram.com/blog/2018/05/27/</a> |
| Node.js FFI     | 18.33 ns  |
| cgo             | 68.77 ns (29x)  |

<https://github.com/dyu/ffi-overhead>



# cgo:

- cmd/cgo
- runtime/cgo
- a sprinkle of cmd/link/internal/ld support

# cgo:

- cmd/cgo
- runtime/cgo
- a sprinkle of cmd/link/internal/ld support
- not a compiler feature!

# cgo:

- `cmd/cgo` — a code generator
- `runtime/cgo`
- a sprinkle of `cmd/link/internal/ld` support
- not a compiler feature!

Reason 1:

calling conventions

C compiler

Go compiler

```
package main

// int my_C_function(int a, int b) {
//     return a + b;
// }
import "C"

func main() {
    a, b := C.int(40), C.int(2)
    c := C.my_C_function(a, b)
    println(a, b, c)
}
```

```
go build -x -work
```

```
// Created by cgo - DO NOT EDIT
```

```
package main
```

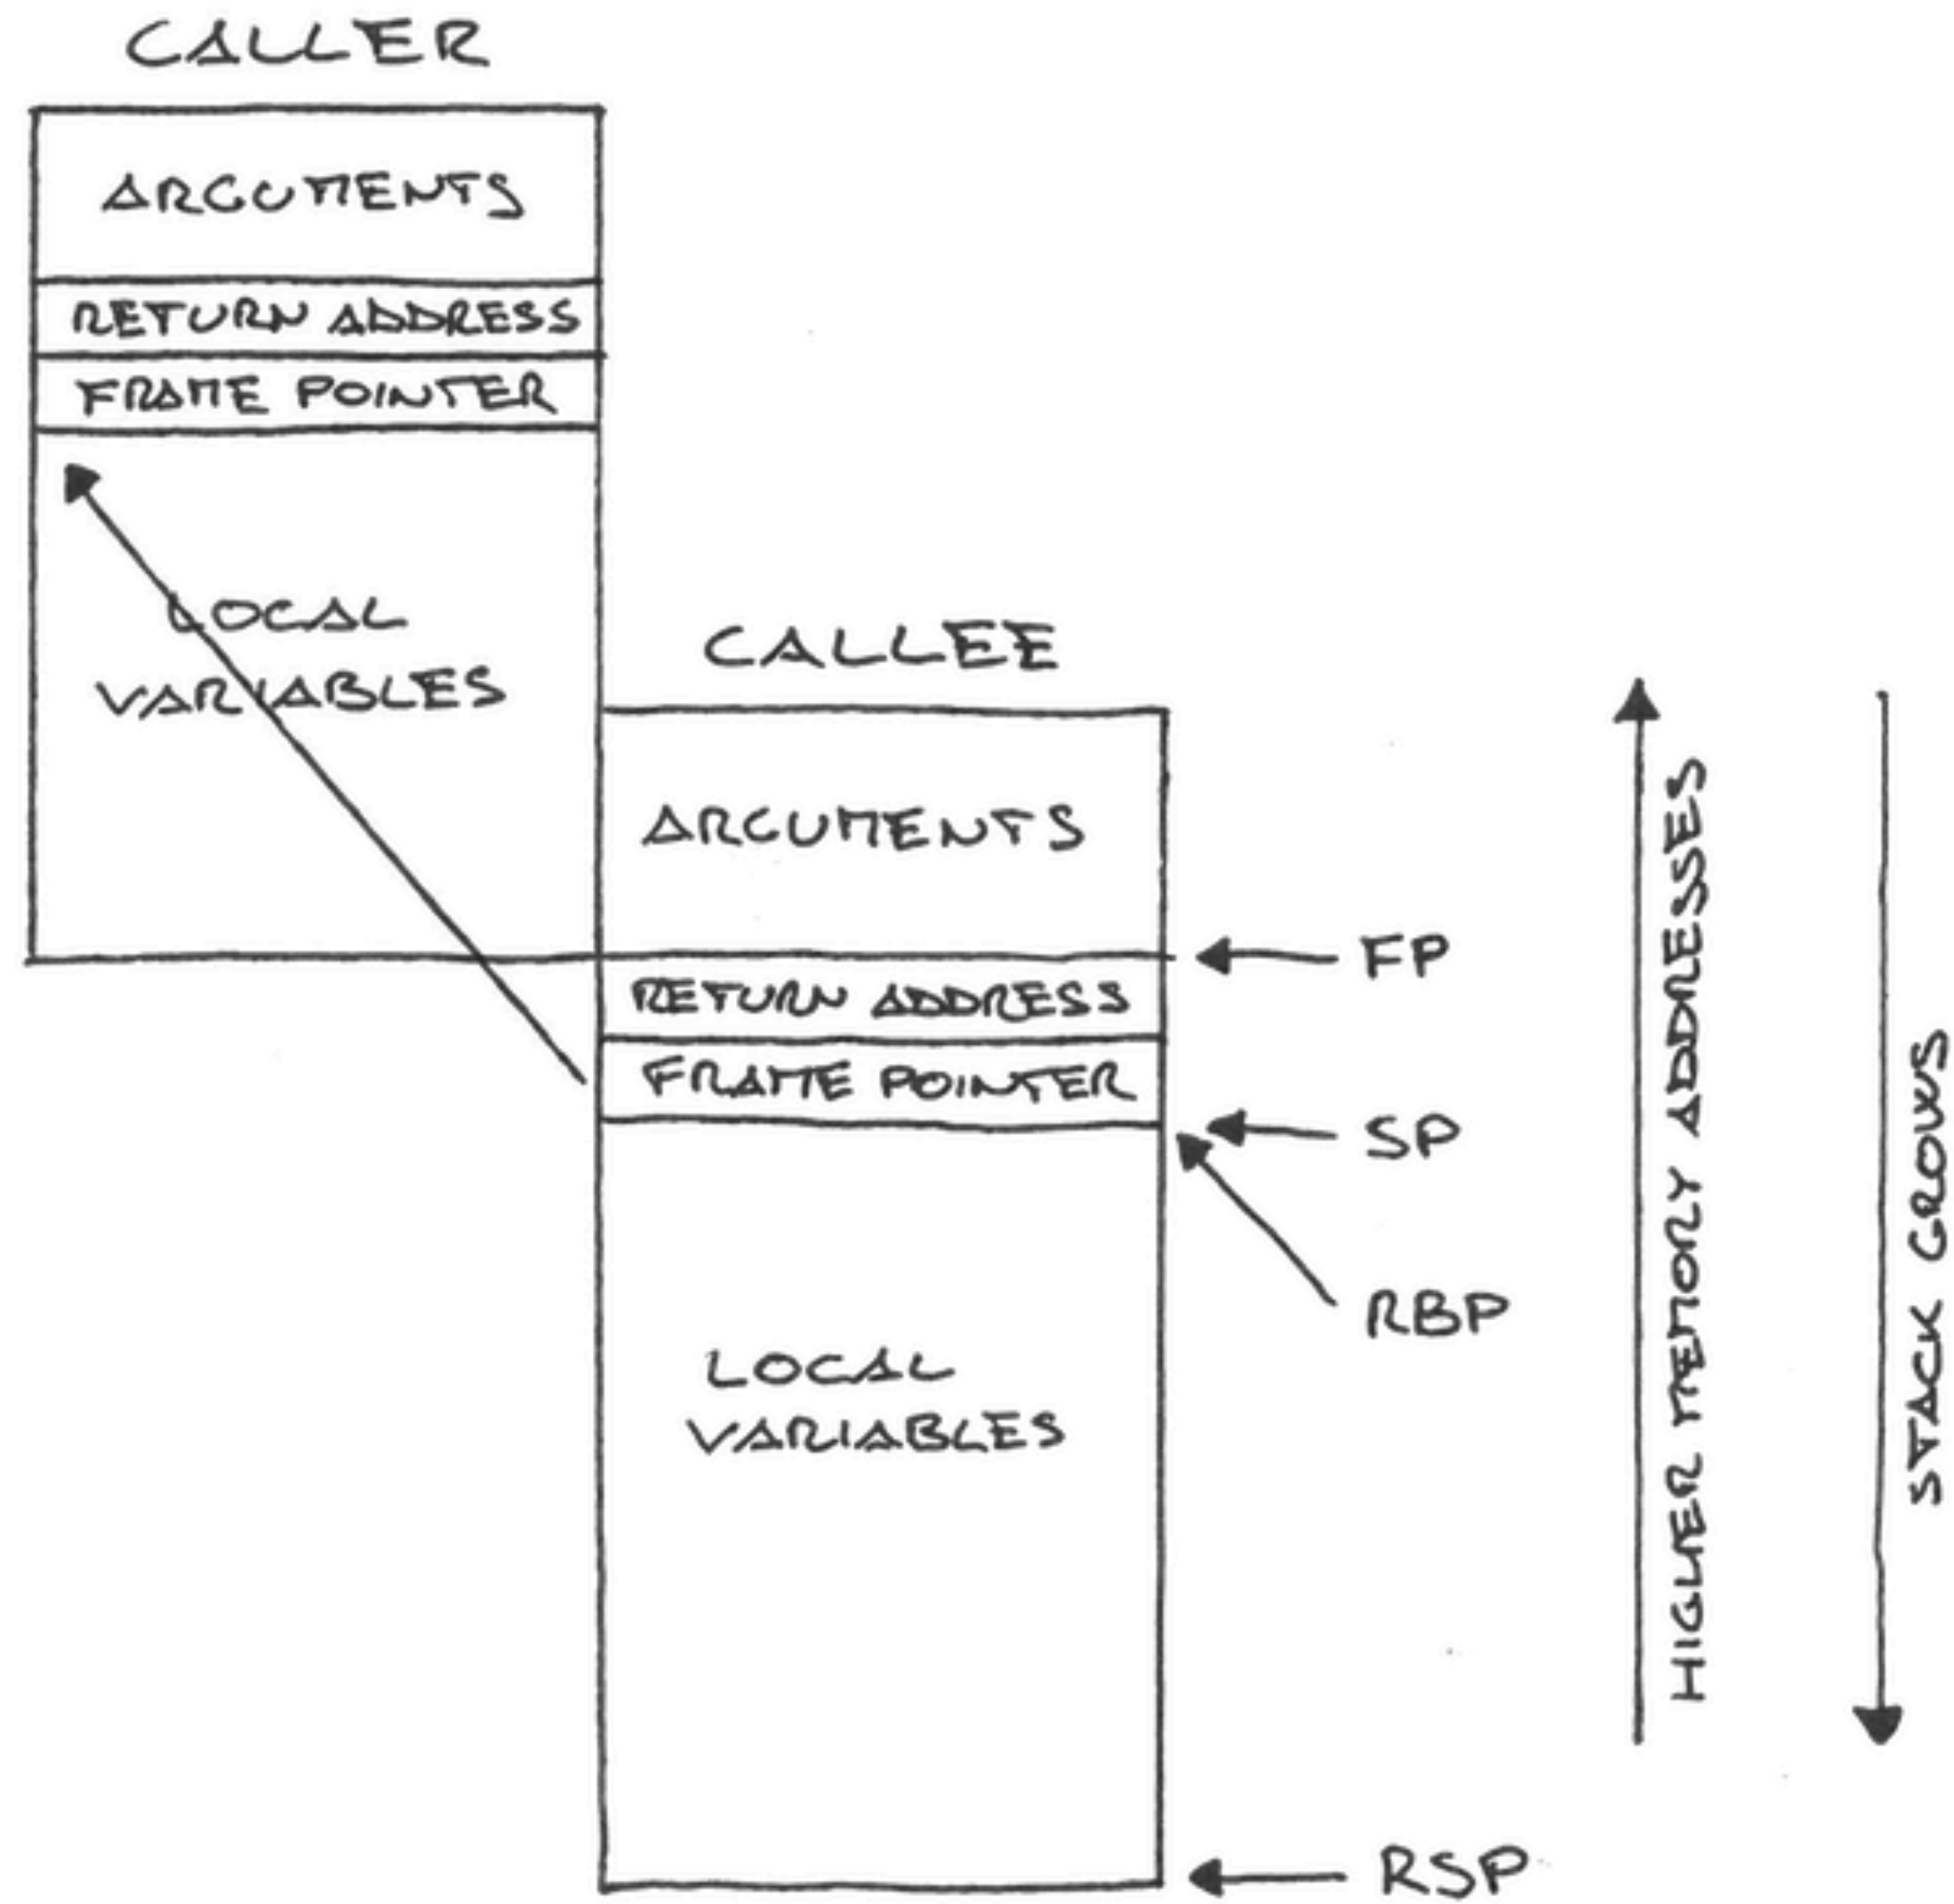
```
// int my_C_function(int a, int b) {  
//     return a + b;  
// }
```

```
import _ "unsafe"
```

```
func main() {  
    a, b := _Ctype_int(40), _Ctype_int(2)  
    c := (_Cfunc_my_C_function)(a, b)  
    println(a, b, c)  
}
```

```
func _Cfunc_my_C_function(p0, p1 _Ctype_int) (r1 _Ctype_int) {  
    runtime.cgocall(_cgo_Cfunc_my_C_function,  
        uintptr(unsafe.Pointer(&p0)))  
    return  
}
```





```
func _Cfunc_my_C_function(p0, p1 _Ctype_int) (r1 _Ctype_int) {  
    runtime.cgocall(_cgo_Cfunc_my_C_function,  
        uintptr(unsafe.Pointer(&p0)))  
    return  
}
```

```
src/runtime/cgocall.go
```

```
// func asmcgocall(fn, arg unsafe.Pointer) int32
// Call fn(arg) on the scheduler stack,
// aligned appropriately for the gcc ABI.
// See cgocall.go for more details.
```

```
TEXT ·asmcgocall(SB),NOSPLIT,$0-20
```

```
    MOVQ    fn+0(FP), AX
```

```
    MOVQ    arg+8(FP), BX
```

```
    MOVQ    BX, DI        // DI = first argument in AMD64 ABI
```

```
    MOVQ    BX, CX        // CX = first argument in Win64
```

```
    CALL    AX
```

```
    MOVL    AX, ret+16(FP)
```

```
    RET
```

```
// func asmcgocall(fn, arg unsafe.Pointer) int32
// Call fn(arg) on the scheduler stack,
// aligned appropriately for the gcc ABI.
// See cgocall.go for more details.
```

```
TEXT ·asmcgocall(SB),NOSPLIT,$0-20
```

```
    MOVQ    fn+0(FP), AX
```

```
    MOVQ    arg+8(FP), BX
```

```
    MOVQ    BX, DI
```

```
    MOVQ    BX, CX
```

```
    CALL    AX
```

```
// DI = first argument in AMD64 ABI
// CX = first argument in Win64
```

```
    MOVL    AX, ret+16(FP)
```

```
    RET
```

```
void _cgo_Cfunc_my_C_function(void *v)
{
    struct {
        int p0;
        int p1;
        int r;
    } *a = v;
    a→r = my_C_function(a→p0, a→p1);
}
```

Go `_Cfunc_my_C_function` →

rewritten Go function

ASM `runtime.[asm]cgocall` →

calling convention  
trampoline

C `_cgo_Cfunc_my_C_function` →

arg unpacking

`my_C_function`

real C function

Learn more:


- `src/runtime/cgocall.go`
- rustgo: Building your own FFI @ GothamGo 2017  
<https://speakerdeck.com/ilosottile/calling-rust-from-go-without-cgo-at-gothamgo-2017>



Reason 2:  
small stacks

Initial goroutine stack size: 2048 bytes


System stack size: 1 – 8 megabytes





stack




function frame
















# Function preamble

```
LEAQ framesize(SP), R0  
CMPQ g→stackguard, R0  
JHI 3(PC)  
MOVQ m→morearg, $(argsize << 32)  
CALL morestack(SB)
```

C doesn't call morestack




C code needs to run on a system stack

```
cgocall / asmcgocall
```


Learn more:

- `src/runtime/stack.go`
- `src/runtime/cgocall.go`
- How stacks are handled in Go by Daniel Morsing  
<https://blog.cloudflare.com/how-stacks-are-handled-in-go/>

Reason 3:  
the scheduler



From <https://morsmachine.dk/go-scheduler>



From <https://morsmachine.dk/go-scheduler>

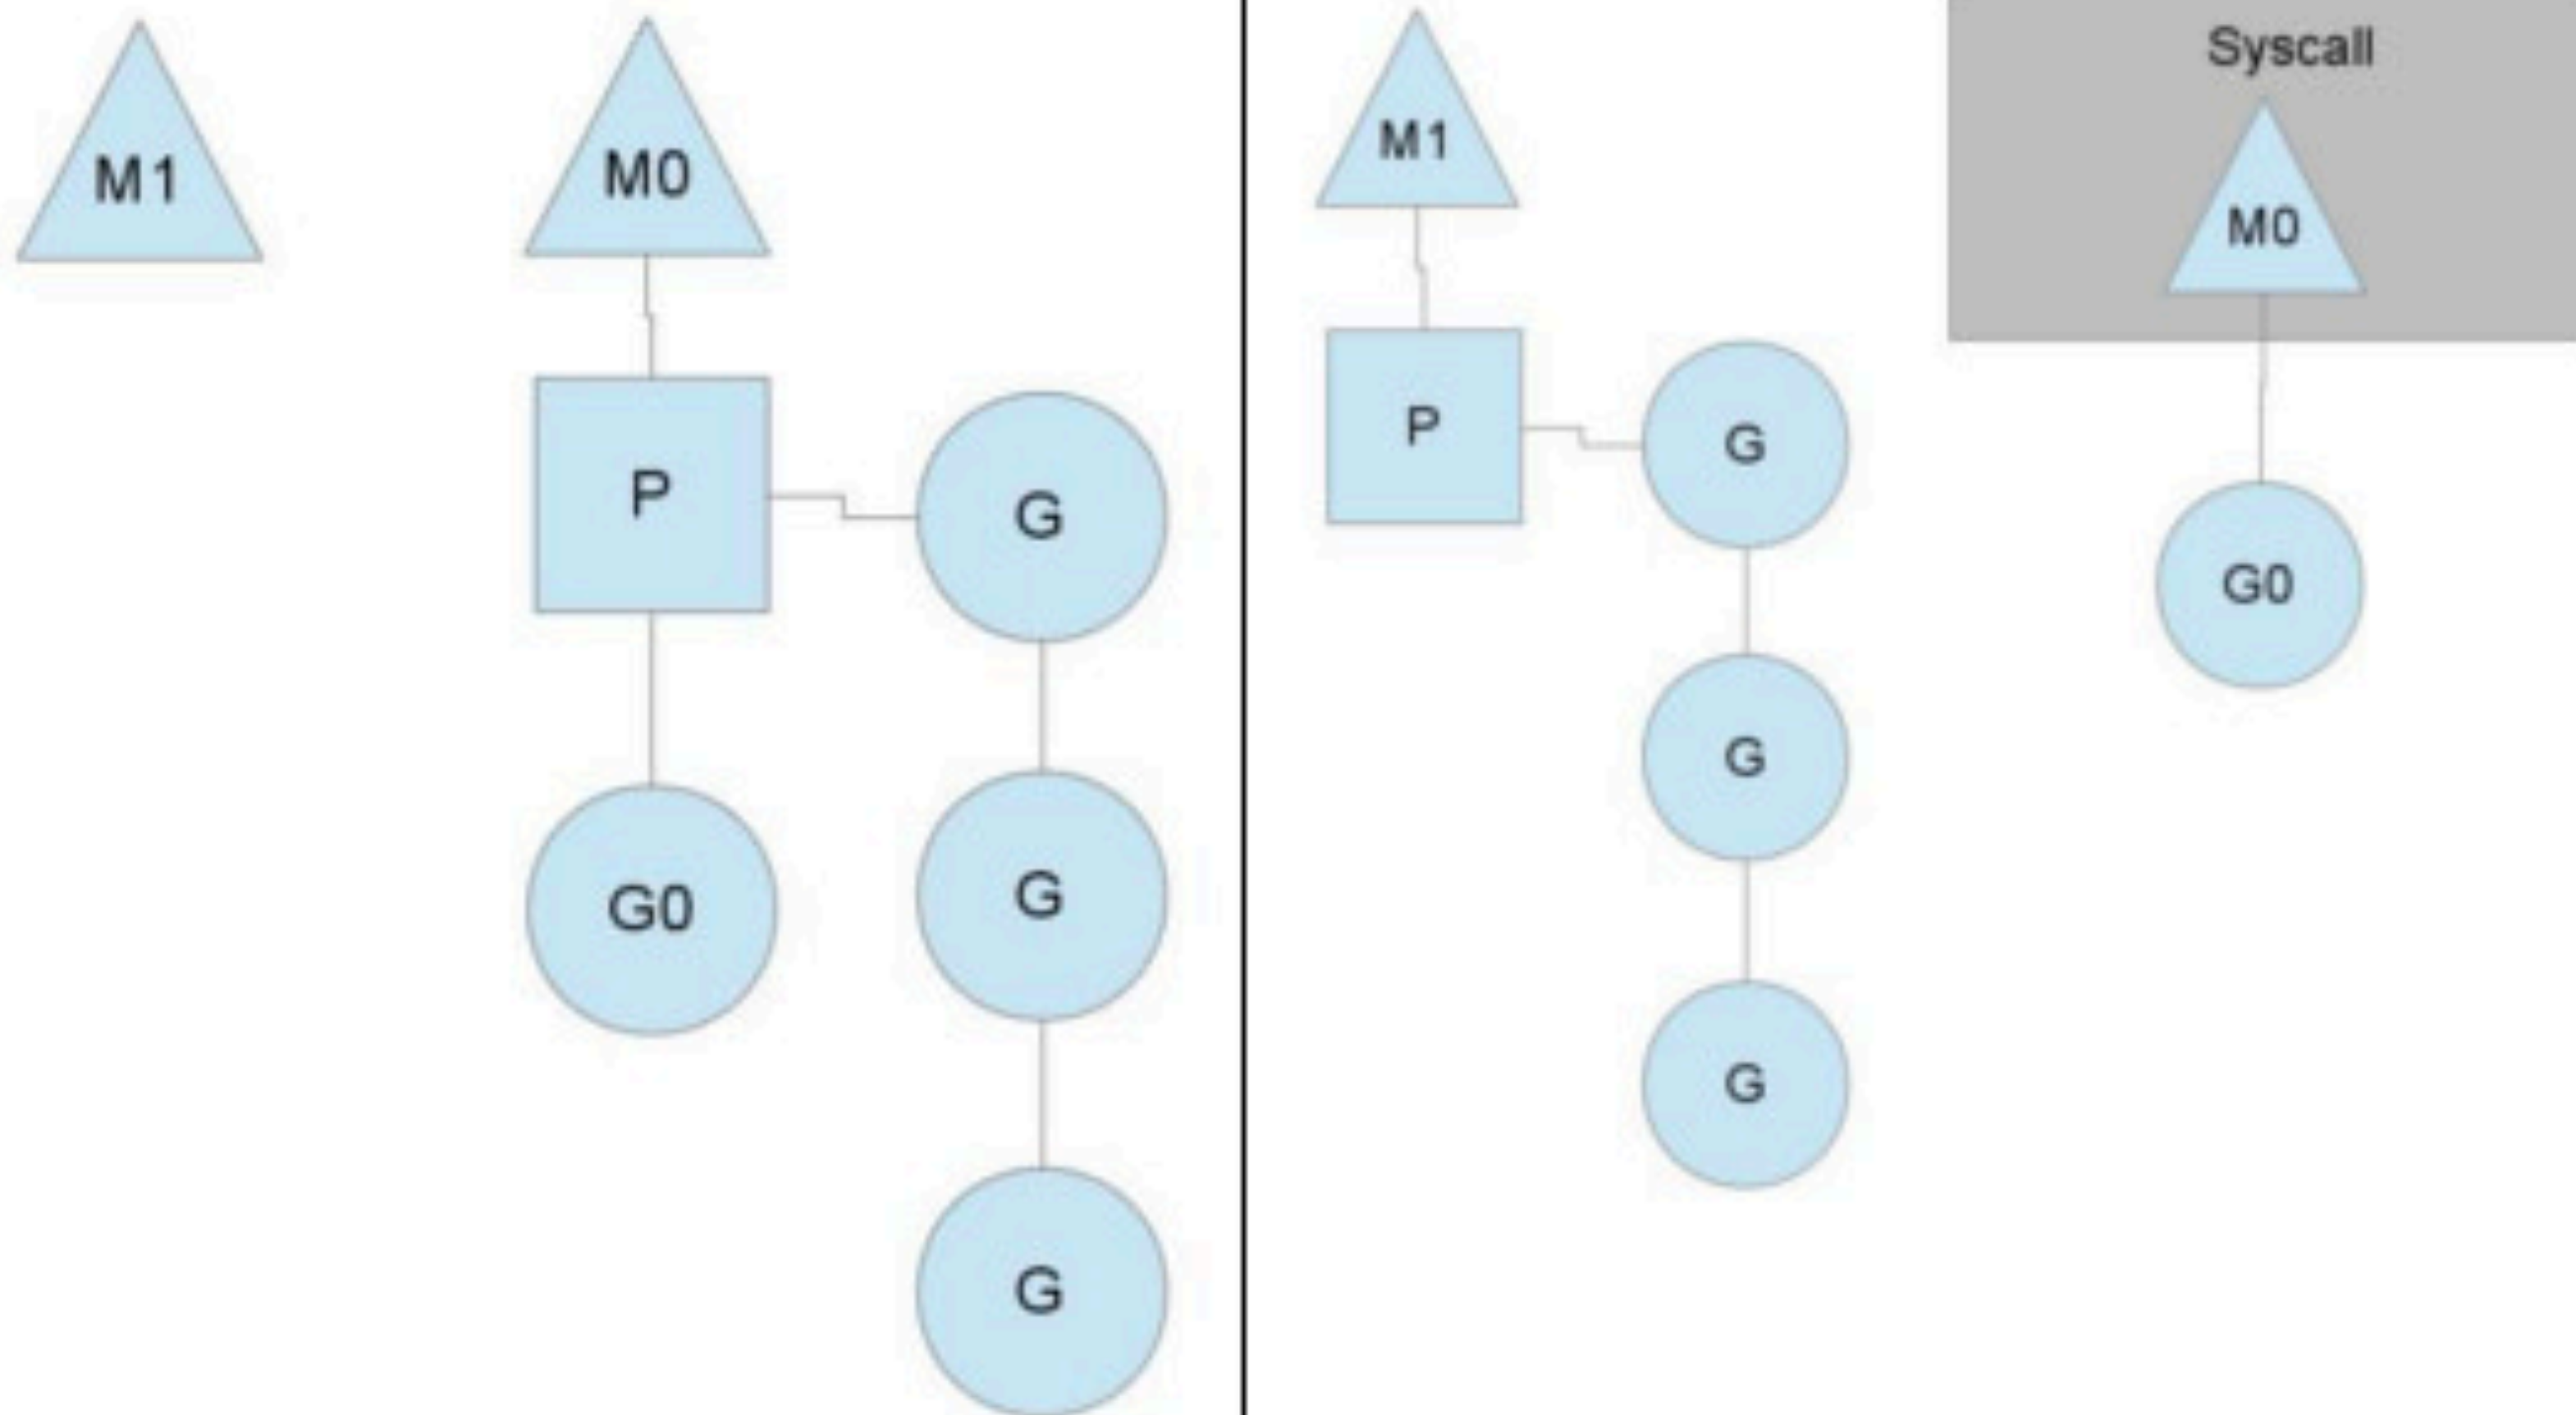
The Go scheduler is collaborative.

It can't preempt running code.

(ProTip: `for {}` is never what you want. Use `select {}`.)



```
▶ // Call from Go to C.  
▶ func cgcoll(fn, arg unsafe.Pointer) int32 {  
    // Announce we are entering a system call  
    // so that the scheduler knows to create another  
    // M to run goroutines while we are in the  
    // foreign code.  
▶     entersyscall()  
  
▶     errno := asmcgcoll(fn, arg)  
  
▶     exitsyscall()
```



From <https://morsmachine.dk/go-scheduler>



Learn more:

- `src/runtime/proc.go` → `reentersyscall`
- The Go scheduler by Daniel Morsing  
<https://morsmachine.dk/go-scheduler>
- Performance without the event loop by Dave Cheney  
<https://dave.cheney.net/2015/08/08/performance-without-the-event-loop>

Reason 4:

the garbage collector

[]byte






Go memory



C memory

[]byte





Go memory



C memory

[]byte




Go memory



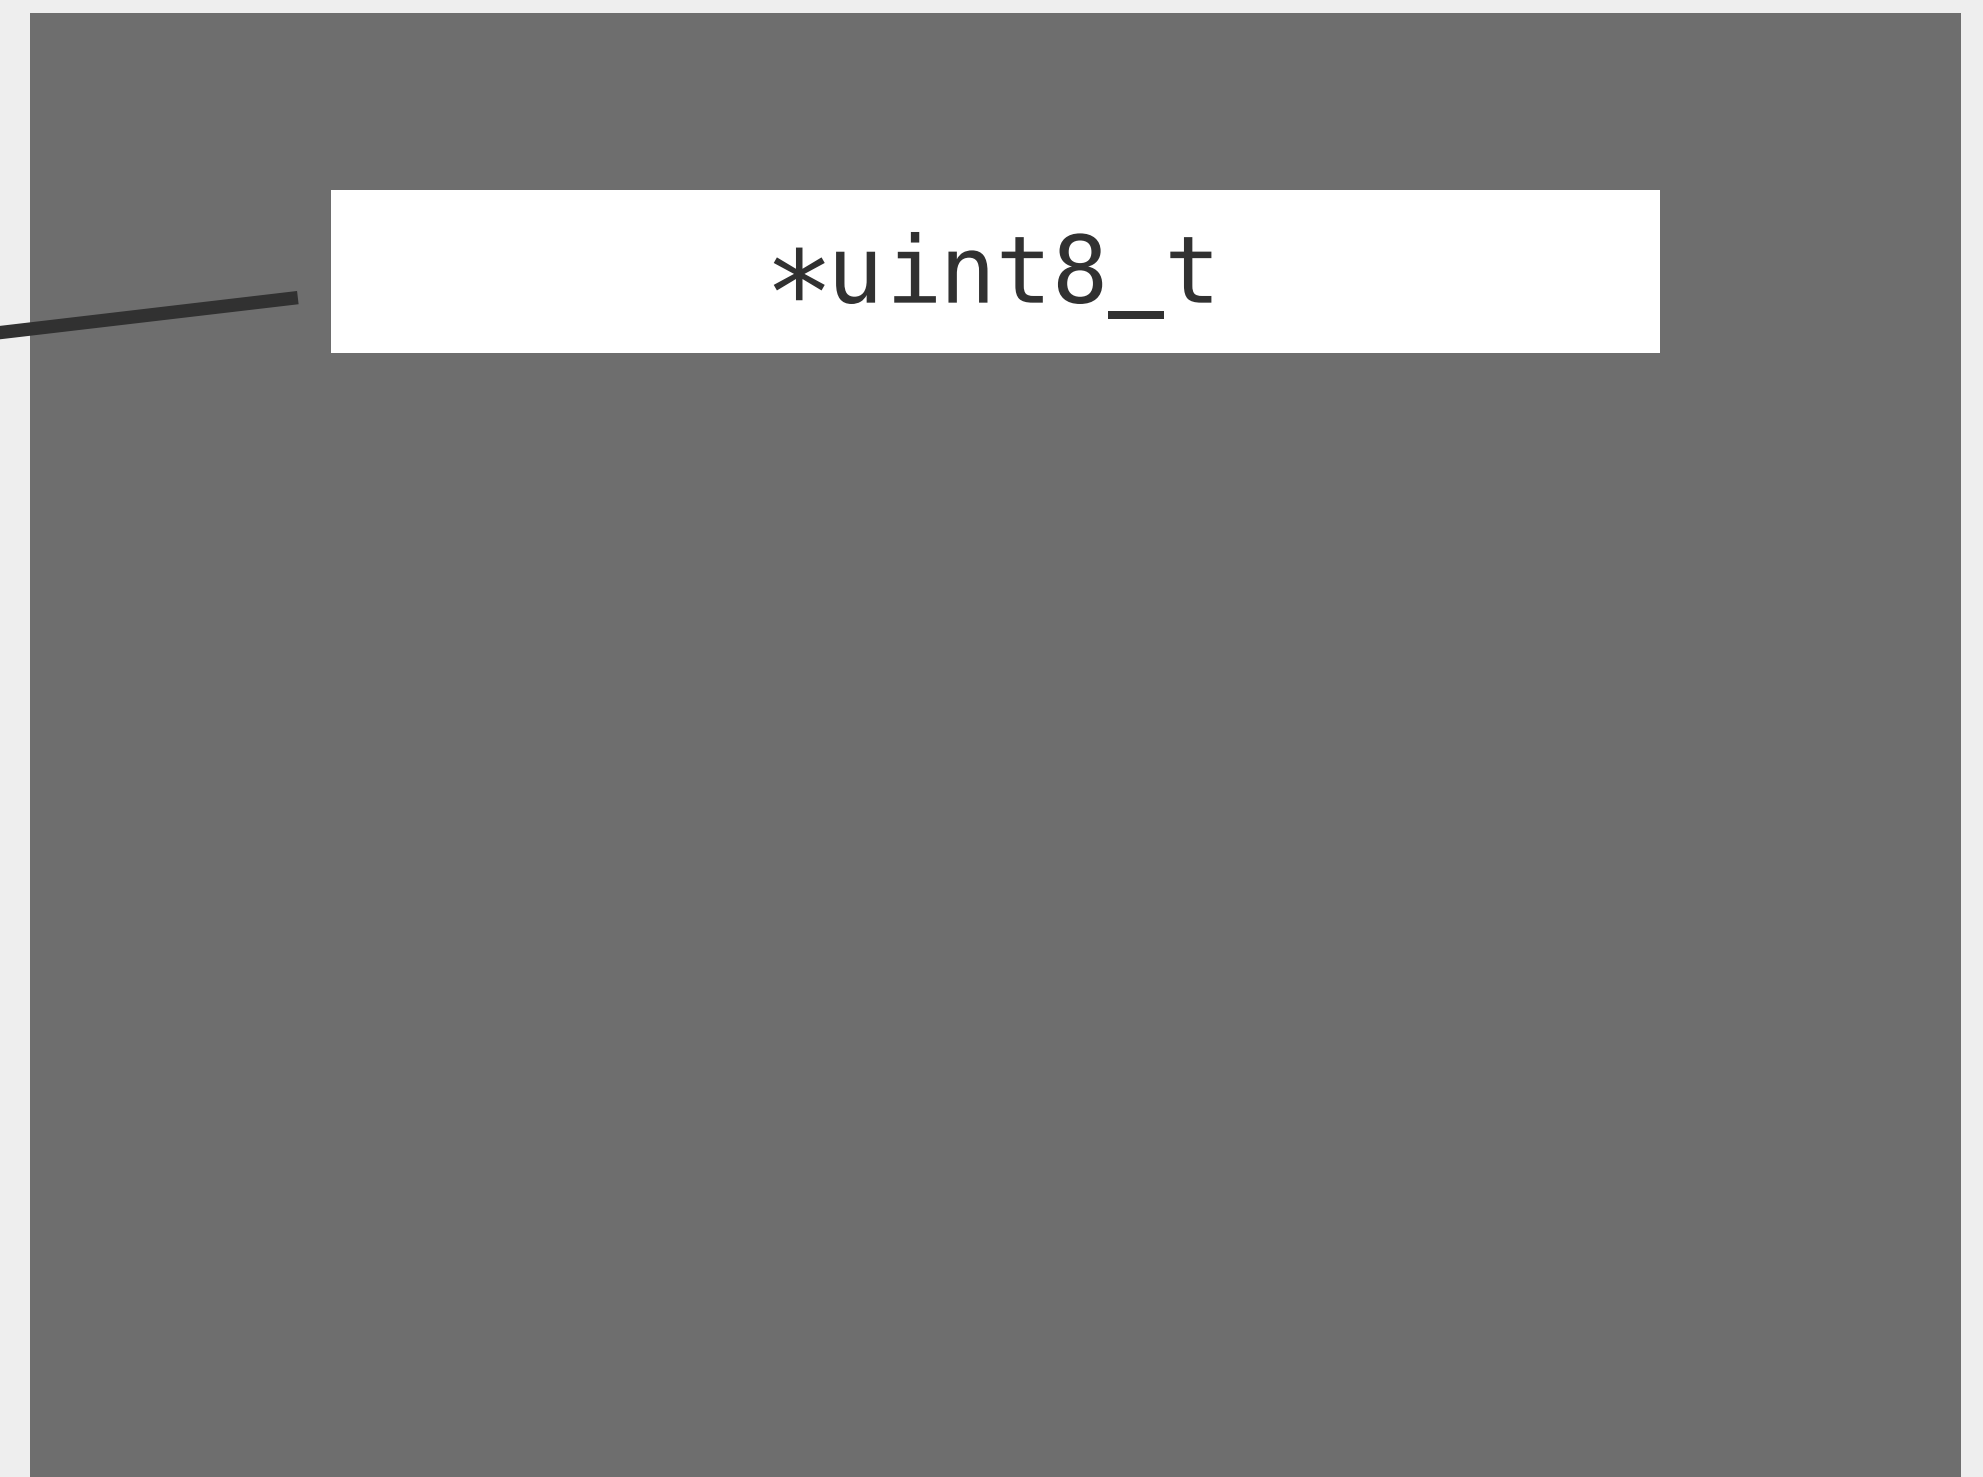
C memory

C.some\_func()

[]byte




Go memory



C memory



# The cgo rules



You may pass a Go pointer

... if it doesn't point to other pointers

... and C can't keep a reference to it



The GC must see all the Go pointers.

```
panic: runtime error: cgo argument  
      has Go pointer to Go pointer
```

```
GODEBUG=cgocheck=2
```

Learn more:

- From cgo back to Go @ GopherCon 2016

<https://speakerdeck.com/ilosottile/from-cgo-back-to-go-gophercon-2016>

# Thank you!

[filippo@golang.org](mailto:filippo@golang.org)

@FiloSottile

